

**Objective:** To reach an accuracy on the test set higher than 45%.

Modify the DNN\_warmup notebook as follows:

- Add suitable hidden layers to the skeleton model in the notebook.
- Add early stopping to the training process in the section titled "Train the model".

```
from tensorflow import keras
from keras.layers import Flatten
from keras.models import Sequential
from keras.layers import Dense
from keras import layers,regularizers
from keras.layers import Dropout,BatchNormalization

tf.keras.backend.clear_session()
np.random.seed(42)
tf.random.set_seed(42)

l2_norm = .0001

# Add a Dense layer with number of neurons equal to the number of classes, with softmax as activation funct
model = Sequential()
model.add(Flatten(input_shape=(32,32,3)))
## you may add here dense layers, using fo instance model.add(Dense(64, activation='relu'))
model.add(Dense(768,activation='relu',kernel_initializer='he_normal',
| | | | kernel_regularizer=regularizers.l2(l2_norm)))
model.add(BatchNormalization())
model.add(Dropout(.2))
model.add(Dense(384,activation='relu',kernel_initializer='he_normal',
| | | | kernel_regularizer=keras.regularizers.l2(l2_norm)))
model.add(BatchNormalization())
model.add(Dropout(.5))
model.add(Dense(num_classes, activation='softmax'))
```

```
from keras.optimizers import Adam, SGD, Adadelta, Adagrad, Adamax, Nadam, RMSprop

sgd = SGD(learning_rate=0.001, momentum=0.09, decay=0.0, nesterov=True)
loss = ['categorical_crossentropy']

metrics = ['accuracy', 'precision', 'recall']
```

✓ 0.1s

Python

```
model.compile(optimizer=sgd,
               loss=loss[0],
               metrics=[metrics[0]],
               )
```

✓ 0.1s

Python

```

batch_size = 128
epochs = 200
from tensorflow import keras
from keras.callbacks import ModelCheckpoint, EarlyStopping
import os

es_callback = EarlyStopping(monitor='val_accuracy', patience=3, verbose=1)

checkpoint_path = "output1/cp.ckpt"
checkpoint_dir = os.path.dirname(checkpoint_path)
cp_callback = ModelCheckpoint(checkpoint_path, monitor='val_accuracy', save_best_only=True)

history = model.fit(x_train, y_train, batch_size=batch_size, validation_data=(x_val, y_val), epochs=epochs,
                    callbacks=[es_callback, cp_callback])

```

✓ 20m 18.1s

Python

Output exceeds the [size limit](#). Open the full output data [in a text editor](#)

Epoch 1/200

293/293 [=====] - ETA: 0s - loss: 3.2264 - accuracy: 0.1911 INFO:tensorflow:Assets written to: output1\cp.ckpt\assets

293/293 [=====] - 14s 45ms/step - loss: 3.2264 - accuracy: 0.1911 - val\_loss: 2.2748 - val\_accuracy: 0.3107

Epoch 2/200

292/293 [=====>.] - ETA: 0s - loss: 2.7906 - accuracy: 0.2574 INFO:tensorflow:Assets written to: output1\cp.ckpt\assets

293/293 [=====] - 13s 44ms/step - loss: 2.7913 - accuracy: 0.2573 - val loss: 2.1151 -

293/293 [=====] - 12s 40ms/step - loss: 1.5343 - accuracy: 0.5373 - val\_loss: 1.5836 - val\_accuracy: 0.5233

Epoch 98/200

293/293 [=====] - 11s 37ms/step - loss: 1.5306 - accuracy: 0.5370 - val\_loss: 1.5813 - val\_accuracy: 0.5232

Epoch 98: early stopping

```

keras.models.load_model('output1\cp.ckpt')

```

✓ 0.7s

Python

<keras.engine.sequential.Sequential at 0x249962e8f70>

## Evaluate the model

```

_, train_acc = model.evaluate(x_train, y_train, verbose=1)
_, val_acc = model.evaluate(x_val, y_val, verbose=1)

_, test_acc = model.evaluate(x_test, y_test, verbose=1)
print('Train: %.3f, val: %.3f, Test: %.3f' % (train_acc, val_acc, test_acc))

```

✓ 16.2s

Python

1172/1172 [=====] - 10s 8ms/step - loss: 1.3386 - accuracy: 0.6242

391/391 [=====] - 3s 8ms/step - loss: 1.5813 - accuracy: 0.5232

313/313 [=====] - 3s 8ms/step - loss: 1.5698 - accuracy: 0.5229

Train: 0.624, val: 0.523, Test: 0.523

**Bonus:** Improve your model to reach an accuracy higher than 50% of the test set.

```
from tensorflow import keras
from keras.layers import Flatten
from keras.models import Sequential
from keras.layers import Dense
from keras import layers, regularizers
from keras.layers import Dropout, Conv2D, Activation, MaxPooling2D, BatchNormalization

tf.keras.backend.clear_session()
np.random.seed(42)
tf.random.set_seed(42)

num_filters_1=32
num_filters_2=64
num_filters_3=128
filter_size=3
pool_size=2
```

```
l2_norm = .0001
model = Sequential()
model.add(Conv2D(num_filters_1, filter_size, padding='same',
                 kernel_regularizer=regularizers.l2(l2_norm), input_shape=x_train.shape[1:], activation='elu'))
model.add(BatchNormalization())

model.add(Conv2D(num_filters_1, filter_size, padding='same',
                 kernel_regularizer=regularizers.l2(l2_norm), activation='elu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=pool_size))
model.add(Dropout(0.2))

model.add(Conv2D(num_filters_2, filter_size, padding='same',
                 kernel_regularizer=regularizers.l2(l2_norm), activation='elu'))
model.add(BatchNormalization())

model.add(Conv2D(num_filters_2, filter_size, padding='same',
                 kernel_regularizer=regularizers.l2(l2_norm), activation='elu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=pool_size))
model.add(Dropout(0.3))

model.add(Conv2D(num_filters_3, filter_size, padding='same',
                 kernel_regularizer=regularizers.l2(l2_norm), activation='elu'))
model.add(BatchNormalization())
model.add(Conv2D(num_filters_3, filter_size, padding='same',
                 kernel_regularizer=regularizers.l2(l2_norm), activation='elu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.4))
```

```
model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))

model.summary()
```

```

from keras.optimizers import Adam, SGD, Adadelta, Adagrad, Adamax, Nadam, RMSprop

rms = RMSprop(learning_rate=0.001, rho=0.9, epsilon=None, decay=0.000001)
# Losses https://keras.io/losses/
loss = ['categorical_crossentropy']

# Metrics https://www.tensorflow.org/api\_docs/python/tf/metrics
metrics = ['accuracy', 'precision', 'recall']

# Compile the model you created before using
# rms optimizer as optimizer
# categorical_crossentropy as loss function
# accuracy as metric
def lr_schedule(epoch):
    lrate = 0.001
    if epoch > 75:
        lrate = 0.0005
    if epoch > 100:
        lrate = 0.0003
    return lrate

model.compile(optimizer=rms,
              loss=loss[0],
              metrics=[metrics[0]],
              )

```

✓ 0.4s

Python

```

batch_size = 64
epochs = 50
from tensorflow import keras
from keras.callbacks import ModelCheckpoint, EarlyStopping
import os

from keras.preprocessing.image import ImageDataGenerator
datagen = ImageDataGenerator( rotation_range=90,
                              width_shift_range=0.1, height_shift_range=0.1,
                              horizontal_flip=True)
datagen.fit(x_train)

es_callback = EarlyStopping(monitor='val_accuracy', patience=10, verbose=1)

checkpoint_path = "output/cp.ckpt"
checkpoint_dir = os.path.dirname(checkpoint_path)
cp_callback = ModelCheckpoint(checkpoint_path, monitor='val_accuracy', save_best_only=True)

history = model.fit(datagen.flow(x_train, y_train, batch_size=batch_size),
                    validation_data=(x_val, y_val), epochs=epochs,
                    callbacks= [keras.callbacks.LearningRateScheduler(lr_schedule), es_callback, cp_callback], verbose=1)

```

✓ 266m 15.6s

Python

```

Epoch 39/50
586/586 [=====] - 138s 235ms/step - loss: 0.9309 - accuracy: 0.7289 - val_loss: 0.9442
- val_accuracy: 0.7394 - lr: 0.0010
Epoch 40/50
586/586 [=====] - 234s 399ms/step - loss: 0.9286 - accuracy: 0.7283 - val_loss: 0.9604
- val_accuracy: 0.7314 - lr: 0.0010
Epoch 41/50
586/586 [=====] - 280s 477ms/step - loss: 0.9254 - accuracy: 0.7311 - val_loss: 0.9528
- val_accuracy: 0.7370 - lr: 0.0010
Epoch 42/50
586/586 [=====] - 276s 471ms/step - loss: 0.9223 - accuracy: 0.7310 - val_loss: 0.8763
- val_accuracy: 0.7566 - lr: 0.0010
Epoch 43/50
586/586 [=====] - 278s 474ms/step - loss: 0.9212 - accuracy: 0.7310 - val_loss: 0.8862
- val_accuracy: 0.7567 - lr: 0.0010
Epoch 43: early stopping

```

```
keras.models.load_model('output\\cp.ckpt')
```

✓ 0.3s

Python

## Evaluate the model

```

_, train_acc = model.evaluate(x_train, y_train, verbose=1)
_, val_acc = model.evaluate(x_val, y_val, verbose=1)

_, test_acc = model.evaluate(x_test, y_test, verbose=1)
print('Train: %.3f, val: %.3f, Test: %.3f' % (train_acc, val_acc, test_acc))

```

✓ 1m 52.7s

Python

```

1172/1172 [=====] - 72s 62ms/step - loss: 0.8200 - accuracy: 0.7703
391/391 [=====] - 19s 49ms/step - loss: 0.8862 - accuracy: 0.7567
313/313 [=====] - 16s 51ms/step - loss: 0.9013 - accuracy: 0.7488
Train: 0.770, val: 0.757, Test: 0.749

```